

RAG (Retrieval-Augmented Generation)

- Split documents into chunks
- Send selected chunks to the LLM instead of entire documents

Sliding Window

- Keep recent conversation messages
- Remove older messages to manage token limits

Summarization

- Compress older conversations or long documents into shorter summaries
- Preserve important information while reducing context usage

Prefix Caching

- Reuse previously computed KV representations for identical prompt prefixes
- Reduce repeated computation, latency, and cost